

Sample Audit Escalation, Evidence & AI Decision-Support Policy

Policy ID	DEMO-AUD-003	Version	1.0
Effective Date	Demo Only	Owner	Audit Operations - Sample
Classification	Synthetic Training Content	Review Cycle	Not Applicable

Important: This document is intentionally fictional and is provided only to test retrieval-augmented generation (RAG), document chunking, search, citations, and audit-assistant workflows. Thresholds and procedures are illustrative and are not payer, coding, regulatory, or clinical guidance.

1. Purpose

This synthetic policy defines a sample governance framework for audit prioritization, escalation, evidence handling, and use of AI-assisted tools. It is designed to support RAG demonstrations in which an AI assistant retrieves internal guidance and cites the relevant section while preserving human decision authority.

2. Core Principles

- Human reviewers retain responsibility for audit conclusions and dispositions.
- Rules and statistical methods generate review signals, not final determinations.
- Generative AI may summarize, retrieve, explain, and organize existing information, but it must not independently approve, deny, or adjudicate a claim.
- Source evidence should be traceable whenever AI-generated content is used in a review workflow.
- Only synthetic or properly de-identified data should be used in the public/demo environment.

3. Priority Queue Framework

Term	Demo Definition
High Priority	Claims with multiple independent signals or a strong rule signal combined with material financial exposure.
Medium Priority	Claims with one or more review signals that warrant additional attention but do not meet the demo high-priority pattern.
Low Priority	Claims with no significant rule or statistical signals under the current demo logic.
Risk Score	A transparent demo score derived from rule and anomaly indicators. It is not a probability of fraud or claim invalidity.

4. Escalation Guidance

Demo Trigger	Indicator	Expected Review	Disposition
Multiple independent signals	Rule signal plus multiple statistical anomalies	Prioritize the claim and gather supporting evidence before escalation.	Priority review
Unresolved material variance	Large financial difference remains unexplained after initial review	Escalate to the designated audit reviewer or lead with documented evidence.	Escalate with evidence
Conflicting or incomplete documentation	Available records do not support a clear reviewer conclusion	Request additional documentation or escalate according to team workflow.	Follow-up / escalation

5. Evidence Requirements

1. Record the claim ID, provider, procedure, billed amount, allowed amount, units, and detected review signals.
2. Capture the specific source document or policy section used to support the reviewer interpretation.
3. Separate retrieved source text from AI-generated explanation in the audit record.
4. Document the human reviewer decision and rationale independently.
5. Retain enough context for another reviewer to reproduce the reasoning from the same evidence.

6. RAG and Generative AI Controls

- RAG retrieval should return relevant policy passages before the LLM generates an explanation.
- The assistant should identify the source document and section used for its answer whenever possible.
- If no relevant source is retrieved, the assistant should state that the knowledge base does not provide enough support rather than inventing a policy.
- AI-generated text should distinguish retrieved facts from interpretation or recommended review steps.
- The assistant must not convert an audit signal into a statement that fraud, overbilling, coding error, medical necessity failure, or claim invalidity has been confirmed.
- A reviewer should be able to inspect the retrieved passages used to generate the response.

7. Data Handling for the Demo

- Use synthetic or properly de-identified claims and policy documents.
- Do not submit real protected health information to free or public AI endpoints without organizational approval and appropriate safeguards.
- API keys and secrets must not be stored in source control.
- Demo output should be treated as decision-support content and reviewed before reuse.

8. Example RAG Questions

- When should a claim be escalated under the sample policy?
- What evidence should be recorded before escalation?
- What is the role of the LLM in the audit workflow?
- What should the assistant do if it cannot retrieve a relevant policy section?
- Can the AI assistant approve or deny a claim?